

# Homework 4

Name: Hanumantha Rao Vakkalanka  
Email: [hvakkala@uncc.edu](mailto:hvakkala@uncc.edu)  
Student ID: 801333188  
Course: Intro to ML

## PROBLEM -1

```
import numpy as np # Importing NumPy for numerical computations
import pandas as pd # Importing Pandas for data manipulation
import matplotlib.pyplot as plt # Importing Matplotlib for data
visualization

from sklearn.preprocessing import StandardScaler # Importing StandardScaler
for feature scaling
from sklearn.preprocessing import MinMaxScaler # Importing MinMaxScaler for
feature scaling
from sklearn.model_selection import KFold # Importing KFold for
cross-validation
from sklearn.model_selection import cross_val_score # Importing
cross_val_score for cross-validation
from sklearn.linear_model import LogisticRegression # Importing
LogisticRegression for classification
from sklearn import datasets # Importing datasets from sklearn
from sklearn import metrics # Importing metrics for performance evaluation
from sklearn.metrics import confusion_matrix # Importing confusion_matrix
for performance evaluation
from sklearn.metrics import classification_report # Importing
classification_report for performance evaluation

from sklearn.datasets import load_breast_cancer # Importing breast cancer
dataset from sklearn
from sklearn.naive_bayes import GaussianNB # Importing GaussianNB for
classification
from sklearn.metrics import precision_recall_curve # Importing
precision_recall_curve for performance evaluation

# Load the breast cancer dataset from sklearn
breast = load_breast_cancer()
```

```

# Extract the features (input data) from the breast cancer dataset
X = breast.data

# Extract the target labels (output data) from the breast cancer dataset
Y = breast.target

# Reshape the feature matrix X to have dimensions (569, 30)
X = np.reshape(X, (569, 30))

# Reshape the target labels Y to have dimensions (569, 1)
Y = np.reshape(Y, (569, 1))

# Create a pandas DataFrame from the reshaped target labels Y
df = pd.DataFrame(Y)

#Doing the SVM for 2 imp components

from sklearn.decomposition import PCA # Import PCA from
sklearn.decomposition

# Initialize PCA with 2 components
pca = PCA(n_components=2)

# Fit PCA to the feature matrix X and transform it to obtain the principal
components
principalComponents = pca.fit_transform(X)

# Create a pandas DataFrame from the principal components with column names
principalDf = pd.DataFrame(data=principalComponents, columns=['principal
component 1', 'principal component 2'])

# Concatenate the principal components DataFrame and the target labels
DataFrame along the columns axis
finalDf1 = pd.concat([principalDf, df], axis=1)

# Display the concatenated DataFrame
finalDf1

```

	principal component 1	principal component 2	0
0	1160.142574	-293.917544	0
1	1269.122443	15.630182	0
2	995.793889	39.156743	0
3	-407.180803	-67.380320	0
4	930.341180	189.340742	0
..	...	...	..
564	1414.126684	110.222492	0
565	1045.018854	77.057589	0
566	314.501756	47.553525	0
567	1124.858115	34.129225	0

568                    -771.527622                    -88.643106    1

[569 rows x 3 columns]

```
# Split the concatenated DataFrame into training and test sets
train = finalDf1.sample(frac=0.8, random_state=0) # Randomly sample 80% for
training set
test = finalDf1.drop(train.index) # Remaining data becomes the test set

# Extract the features (principal component 1) and target labels from the
training and test sets
X_train = train.values[:, 0:1] # Extract the first column (principal
component 1) as X_train
X_test = test.values[:, 0:1] # Extract the first column (principal component
1) as X_test

# Extract the target labels (column 2) from the training and test sets
Y_train = train.values[:, 2] # Extract column 2 (target labels) as Y_train
Y_test = test.values[:, 2] # Extract column 2 (target labels) as Y_test

# Display the extracted target labels from the test set
Y_test
array([0., 0., 0., 0., 0., 0., 0., 0., 1., 1., 0., 0., 0., 1., 0., 0., 0.,
       1., 0., 0., 0., 0., 1., 0., 1., 0., 1., 1., 1., 0., 1., 1., 1., 1.,
       1., 1., 1., 0., 1., 0., 1., 1., 1., 0., 0., 0., 0., 1., 1., 1., 0.,
       0., 0., 0., 1., 1., 0., 1., 1., 1., 1., 1., 1., 1., 0., 1., 0., 1.,
       1., 1., 0., 0., 0., 0., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,
       1., 1., 0., 1., 1., 1., 0., 1., 0., 1., 1., 1., 1., 0., 1., 1., 0.,
       1., 1., 1., 0., 1., 1., 1., 1., 1., 1., 1., 1.])

# Import GridSearchCV and SVC from sklearn
from sklearn.model_selection import GridSearchCV
from sklearn.svm import SVC

# Define the hyperparameter grid for GridSearchCV
param_grid = {'C': [0.1, 1, 10, 100, 1000],
              'gamma': [1, 0.1, 0.01, 0.001, 0.0001],
              'kernel': ['rbf']}

# Create an instance of GridSearchCV with SVC as the estimator
grid = GridSearchCV(SVC(), param_grid, refit=True, cv=5)

# Fit the GridSearchCV to the training data
grid.fit(X_train, Y_train)

GridSearchCV(cv=5, estimator=SVC(),
             param_grid={'C': [0.1, 1, 10, 100, 1000],
                        'gamma': [1, 0.1, 0.01, 0.001, 0.0001],
                        'kernel': ['rbf']})
```

```
# Fit the GridSearchCV to the training data
```

```
grid.fit(X_train, Y_train)
```

```
# Predict the target labels on the test data
```

```
Y_predicted = grid.predict(X_test)
```

```
# Display the classification report
```

```
print(metrics.classification_report(Y_test, Y_predicted))
```

```
# Display the confusion matrix
```

```
print(metrics.confusion_matrix(Y_test, Y_predicted))
```

	precision	recall	f1-score	support
0.0	0.88	0.82	0.85	44
1.0	0.89	0.93	0.91	70
accuracy			0.89	114
macro avg	0.88	0.87	0.88	114
weighted avg	0.89	0.89	0.89	114

```
[[36  8]  
 [ 5 65]]
```

```
# Calculate precision and recall
```

```
precision, recall, thresholds = precision_recall_curve(Y_test, Y_predicted)
```

```
# Create precision-recall curve
```

```
fig, ax = plt.subplots()
```

```
ax.plot(recall, precision, color='purple')
```

```
# Add axis labels to the plot
```


```
ax.set_title('Precision-Recall Curve')
```

```
ax.set_ylabel('Precision')
```

```
ax.set_xlabel('Recall')
```

```
# Display the plot
```

```
plt.show()
```



*# Doing SVM for 6 important components*

```
from sklearn.decomposition import PCA
```

*# Perform PCA with 6 components*

```
pca = PCA(n_components=6)
principalComponents = pca.fit_transform(X)
principalDf = pd.DataFrame(data=principalComponents,
                           columns=['principal component 1', 'principal
principal component 2', 'principal component 3',
                                   'principal component 4', 'principal
principal component 5', 'principal component 6'])
```

*# Concatenate PCA components with target variable*

```
finalDf2 = pd.concat([principalDf, df], axis=1)
```

	principal component 1	principal component 2	principal component 3 \
0	1160.142574	-293.917544	48.578398
1	1269.122443	15.630182	-35.394534
2	995.793889	39.156743	-1.709753
3	-407.180803	-67.380320	8.672848
4	930.341180	189.340742	1.374801
..	...	...	...
564	1414.126684	110.222492	40.065944
565	1045.018854	77.057589	0.036669
566	314.501756	47.553525	-10.442407
567	1124.858115	34.129225	-19.742087

568	-771.527622	-88.643106	23.889032	
	principal component 4	principal component 5	principal component 6	0
0	-8.711975	32.000486	1.265415	0
1	17.861283	-4.334874	-0.225872	0
2	4.199340	-0.466529	-2.652811	0
3	-11.759867	7.115461	1.299436	0
4	8.499183	7.613289	1.021160	0
..	...	...	...	..
564	6.562240	-5.102856	-0.395424	0
565	-4.753245	-12.417863	-0.059637	0
566	-9.771881	-6.156213	-0.870726	0
567	-23.660881	3.565133	4.086390	0
568	2.547249	-14.717566	4.418123	1

```
[569 rows x 7 columns]
```

## # Splitting data into train and test sets

```
train2 = finalDf2.sample(frac=0.8, random_state=0)
test2 = finalDf2.drop(train2.index)
```

### # Extracting features and target variables from train and test sets

```
X_train2 = train2.values[:, 0:5]
X_test2 = test2.values[:, 0:5]
```

```
Y_train2 = train2.values[:, 6]
Y_test2 = test2.values[:, 6]
```

```
array([0., 0., 0., 0., 0., 0., 0., 0., 1., 1., 0., 0., 0., 1., 0., 0., 0.,
       1., 0., 0., 0., 0., 1., 0., 1., 0., 1., 1., 1., 0., 1., 1., 1., 1.,
       1., 1., 1., 0., 1., 0., 1., 1., 1., 1., 0., 0., 0., 0., 1., 1., 1., 0.,
       0., 0., 0., 1., 1., 0., 1., 1., 1., 1., 1., 1., 1., 0., 1., 0., 1.,
       1., 1., 0., 0., 0., 0., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,
       1., 1., 0., 1., 1., 1., 0., 1., 0., 1., 1., 1., 1., 0., 1., 1., 0.,
       1., 1., 1., 0., 1., 1., 1., 1., 1., 1., 1., 1., 1.] )
```

### # Performing Grid Search with SVM on the reduced-dimensionality data

```
from sklearn.model_selection import GridSearchCV
from sklearn.svm import SVC
```

```
# Define the hyperparameter grid for Grid Search
```

```
param_grid = {'C': [0.1, 1, 10, 100, 1000],
              'gamma': [1, 0.1, 0.01, 0.001, 0.0001],
              'kernel': ['rbf']}
```

```
# Create a GridSearchCV object with SVM classifier, hyperparameter grid, and 5-fold cross-validation
```

```

grid = GridSearchCV(SVC(), param_grid, refit=True, cv=5)

# Fit the grid search object to the reduced-dimensionality training data
grid.fit(X_train2, Y_train2)

GridSearchCV(cv=5, estimator=SVC(),
             param_grid={'C': [0.1, 1, 10, 100, 1000],
                        'gamma': [1, 0.1, 0.01, 0.001, 0.0001],
                        'kernel': ['rbf']})

# Fitting the GridSearchCV object to the reduced-dimensionality training data

grid.fit(X_train2, Y_train2)

GridSearchCV(cv=5, estimator=SVC(),
             param_grid={'C': [0.1, 1, 10, 100, 1000],
                        'gamma': [1, 0.1, 0.01, 0.001, 0.0001],
                        'kernel': ['rbf']})

# Making predictions on the reduced-dimensionality test data

Y_predicted = grid.predict(X_test2)
Y_predicted

array([0., 0., 0., 0., 1., 0., 0., 0., 1., 1., 0., 0., 0., 1., 0., 0., 0.,
       1., 0., 1., 0., 0., 1., 1., 1., 0., 1., 1., 1., 0., 1., 1., 1., 1.,
       1., 1., 1., 1., 1., 0., 1., 1., 1., 1., 1., 0., 0., 0., 1., 1., 1., 0.,
       0., 0., 0., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 0., 1., 0., 1.,
       1., 1., 0., 0., 0., 0., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,
       1., 1., 1., 1., 1., 1., 0., 1., 0., 1., 1., 1., 1., 0., 1., 1., 0.,
       1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.])

# Printing classification report and confusion matrix

from sklearn import metrics

print(metrics.classification_report(Y_test2, Y_predicted))
print(metrics.confusion_matrix(Y_test2, Y_predicted))

# Calculating precision and recall

precision, recall, thresholds = precision_recall_curve(Y_test2, Y_predicted)

# Creating precision-recall curve plot

fig, ax = plt.subplots()
ax.plot(recall, precision, color='purple')

# Adding axis labels to the plot

```


```
ax.set_title('Precision-Recall Curve')
ax.set_ylabel('Precision')
ax.set_xlabel('Recall')
```

```
# Displaying the plot
```

```
plt.show()
```

	precision	recall	f1-score	support
0.0	1.00	0.82	0.90	44
1.0	0.90	1.00	0.95	70
accuracy			0.93	114
macro avg	0.95	0.91	0.92	114
weighted avg	0.94	0.93	0.93	114

```
[[36  8]
 [ 0 70]]
```



```
# Performing SVM for 18 important components
```

```
from sklearn.decomposition import PCA
```

```
# Applying PCA with n_components = 18
```

```
pca = PCA(n_components=18)
principalComponents = pca.fit_transform(X)
```



```
# Creating a DataFrame with the principal components as columns
principalDf = pd.DataFrame(data = principalComponents,
                           columns = ['principal component 1', 'principal
component 2','principal component 3', 'principal component 4', 'principal
component 5','principal component 6', 'principal component 7','principal
component 8', 'principal component 9', 'principal component 10', 'principal
component 11','principal component 12', 'principal component 13','principal
component 14', 'principal component 15', 'principal component 16','principal
component 17', 'principal component 18'])
```

```
# Concatenating the principal components DataFrame with the original
DataFrame along the columns axis
finalDf3 = pd.concat([principalDf, df], axis = 1)
finalDf3
```

	principal component 1	principal component 2	principal component 3	\
0	1160.142574	-293.917544	48.578398	
1	1269.122443	15.630182	-35.394534	
2	995.793889	39.156743	-1.709753	
3	-407.180803	-67.380320	8.672848	
4	930.341180	189.340742	1.374801	
..	...	...	...	
564	1414.126684	110.222492	40.065944	
565	1045.018854	77.057589	0.036669	
566	314.501756	47.553525	-10.442407	
567	1124.858115	34.129225	-19.742087	
568	-771.527622	-88.643106	23.889032	

	principal component 4	principal component 5	principal component 6	\
0	-8.711975	32.000486	1.265415	
1	17.861283	-4.334874	-0.225872	
2	4.199340	-0.466529	-2.652811	
3	-11.759867	7.115461	1.299436	
4	8.499183	7.613289	1.021160	
..	...	...	...	
564	6.562240	-5.102856	-0.395424	
565	-4.753245	-12.417863	-0.059637	
566	-9.771881	-6.156213	-0.870726	
567	-23.660881	3.565133	4.086390	
568	2.547249	-14.717566	4.418123	

	principal component 7	principal component 8	principal component 9	\
0	0.931337	0.148167	0.745463	
1	-0.046037	0.200804	-0.485828	

2	-0.779745	-0.274026	-0.173874
3	-1.267304	-0.060555	-0.330639
4	-0.335522	0.289109	0.036087
..	...	...	...
564	-0.786751	0.037082	-0.452530
565	0.449831	0.509154	-0.449986
566	-2.166493	-0.442279	-0.097398
567	-1.705401	-0.359964	0.385030
568	-2.815752	0.030039	-0.423451

	principal component 10	principal component 11	principal component 12
\			
0	0.589359	-0.307804	0.043452
1	-0.084035	0.080642	0.033042
2	-0.186994	0.279174	-0.020464
3	-0.144155	0.927471	-0.174720
4	-0.138502	0.042228	-0.062721
..	...	...	...
564	-0.235185	0.163649	0.052543
565	0.493247	0.007625	0.055832
566	-0.144667	-0.109147	0.076263
567	0.615467	0.307166	-0.028224
568	-0.301439	0.133353	-0.115105

	principal component 13	principal component 14	principal component 15
\			
0	0.034777	0.065069	-0.012934
1	0.045485	-0.005534	0.021368
2	0.083505	0.024824	-0.026887
3	0.282556	0.080057	0.043201
4	-0.114247	0.002274	-0.019548
..	...	...	...
564	-0.075032	-0.015211	-0.061390
565	-0.015163	0.009985	0.003312
566	-0.004448	-0.055285	-0.012459
567	0.060561	-0.037742	-0.031873
568	-0.019667	0.013734	-0.004134

	principal component 16	principal component 17	principal component 18
0			
0	-0.002670	0.018300	0.010263
0			
1	-0.028715	0.012371	-0.006009
0			
2	-0.041255	0.008218	-0.028044
0			
3	-0.034175	0.033742	-0.016965
0			
4	0.019932	-0.019201	0.004024
0			

```

..          ...          ...          ...
..
564          -0.054694          -0.004829          -0.011515
0
565          -0.020654          0.005197          0.002106
0
566          -0.005414          0.007866          -0.004484
0
567          0.020126          0.015243          0.043651
0
568          0.034264          0.009440          -0.028323
1

```

```
[569 rows x 19 columns]
```

```
# Randomly sample 80% of rows from finalDf3 to create the training dataset
```

```
# Setting random_state to 0 for reproducibility
```

```
train3 = finalDf3.sample(frac=0.8, random_state=0)
```

```
# Drop the rows that were sampled for the training dataset to create the test dataset
```

```
test3 = finalDf3.drop(train3.index)
```

```
# Extract the feature columns (principal components) from the training dataset
```

```
X_train3 = train3.values[:, 0:17]
```

```
# Extract the feature columns (principal components) from the test dataset
```

```
X_test3 = test3.values[:, 0:17]
```

```
# Extract the target column from the training dataset
```

```
Y_train3 = train3.values[:, 18]
```

```
# Extract the target column from the test dataset
```

```
Y_test3 = test3.values[:, 18]
```

```
array([0., 0., 0., 0., 0., 0., 0., 0., 1., 1., 0., 0., 0., 1., 0., 0., 0.,
       1., 0., 0., 0., 0., 1., 0., 1., 0., 1., 1., 1., 0., 1., 1., 1., 1.,
       1., 1., 1., 0., 1., 0., 1., 1., 1., 0., 0., 0., 0., 1., 1., 1., 0.,
       0., 0., 0., 1., 1., 0., 1., 1., 1., 1., 1., 1., 1., 0., 1., 0., 1.,
       1., 1., 0., 0., 0., 0., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,
       1., 1., 0., 1., 1., 1., 0., 1., 0., 1., 1., 1., 1., 0., 1., 1., 0.,
       1., 1., 1., 0., 1., 1., 1., 1., 1., 1., 1., 1.])
```

```
# Import necessary libraries for GridSearchCV and SVM
```

```
from sklearn.model_selection import GridSearchCV
```

```
from sklearn.svm import SVC
```

```
# Define the hyperparameter grid for the SVM model
```

```
param_grid = {'C': [0.1, 1, 10, 100, 1000],
```

```

        'gamma': [1, 0.1, 0.01, 0.001, 0.0001],
        'kernel': ['rbf']}]

# Create an instance of GridSearchCV with SVC as the base estimator, the
# defined parameter grid, refit=True for model re-fitting, and cv=5 for 5-fold
# cross-validation
grid = GridSearchCV(SVC(), param_grid, refit=True, cv=5)

# Fit the GridSearchCV object to the training data
grid.fit(X_train3, Y_train3)

GridSearchCV(cv=5, estimator=SVC(),
             param_grid={'C': [0.1, 1, 10, 100, 1000],
                         'gamma': [1, 0.1, 0.01, 0.001, 0.0001],
                         'kernel': ['rbf']})

# Fit the GridSearchCV object to the training data
grid.fit(X_train3, Y_train3)

GridSearchCV(cv=5, estimator=SVC(),
             param_grid={'C': [0.1, 1, 10, 100, 1000],
                         'gamma': [1, 0.1, 0.01, 0.001, 0.0001],
                         'kernel': ['rbf']})

Y_predicted = grid.predict(X_test3)
Y_predicted

array([0., 0., 0., 0., 1., 0., 0., 0., 1., 1., 0., 0., 0., 1., 0., 0., 0.,
       1., 0., 1., 0., 0., 1., 1., 1., 0., 1., 1., 1., 0., 1., 1., 1., 1.,
       1., 1., 1., 1., 1., 0., 1., 1., 1., 1., 1., 0., 0., 0., 1., 1., 1., 0.,
       0., 0., 0., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 0., 1., 0., 1.,
       1., 1., 0., 0., 0., 0., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,
       1., 1., 1., 1., 1., 1., 0., 1., 0., 1., 1., 1., 1., 0., 1., 1., 0.,
       1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.])

print(metrics.classification_report(Y_test, Y_predicted))

print(metrics.confusion_matrix(Y_test, Y_predicted))

#calculate precision and recall
precision, recall, thresholds = precision_recall_curve(Y_test, Y_predicted)
#create precision recall curve
fig, ax = plt.subplots()
ax.plot(recall, precision, color='purple')


#add axis labels to plot
ax.set_title('Precision-Recall Curve')
ax.set_ylabel('Precision')
ax.set_xlabel('Recall')

```

```
#display plot
plt.show()
```

	precision	recall	f1-score	support
0.0	1.00	0.82	0.90	44
1.0	0.90	1.00	0.95	70
accuracy			0.93	114
macro avg	0.95	0.91	0.92	114
weighted avg	0.94	0.93	0.93	114

```
[[36  8]
 [ 0 70]]
```



```
# Trying different kernels
```

```
# SVM with linear kernel
```

```
# SVM with linear kernel and C=100.0
```

```
linear_svc=SVC(kernel='linear', C=100.0)
```

```
linear_svc.fit(X_train,Y_train)
```

```
# make predictions on test set
```

```
Y_pred_test=linear_svc.predict(X_test)
```

```
# compute and print accuracy score
```

```
print(metrics.classification_report(Y_test, Y_pred_test))
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0.0	0.95	0.84	0.89	44
1.0	0.91	0.97	0.94	70
accuracy			0.92	114
macro avg	0.93	0.91	0.91	114
weighted avg	0.92	0.92	0.92	114

```
#SVM with polynomial kernel and C=1000.0
poly_svc100=SVC(kernel='poly', C=1000.0)
# fit classifier to training set
poly_svc100.fit(X_train, Y_train)
# make predictions on test set
Y_pred=poly_svc100.predict(X_test)
# compute and print accuracy score
print(metrics.classification_report(Y_test, Y_pred))
```

	precision	recall	f1-score	support
0.0	1.00	0.66	0.79	44
1.0	0.82	1.00	0.90	70
accuracy			0.87	114
macro avg	0.91	0.83	0.85	114
weighted avg	0.89	0.87	0.86	114

```
# Run SVM with sigmoid kernel
# instantiate classifier with sigmoid kernel and C=100.0
sigmoid_svc100=SVC(kernel='sigmoid', C=100.0)
# fit classifier to training set
sigmoid_svc100.fit(X_train,Y_train)
# make predictions on test set
Y_pred=sigmoid_svc100.predict(X_test)
# compute and print accuracy score
print(metrics.classification_report(Y_test, Y_pred))
```

	precision	recall	f1-score	support
0.0	0.75	0.75	0.75	44
1.0	0.84	0.84	0.84	70
accuracy			0.81	114
macro avg	0.80	0.80	0.80	114
weighted avg	0.81	0.81	0.81	114

```
#accuracy graph
count = 30
while(count >= 2):
```



```
0.9298245614035088
0.9298245614035088
0.9298245614035088
0.9298245614035088
0.9298245614035088
0.9298245614035088
0.9298245614035088
0.9298245614035088
0.9298245614035088
0.9298245614035088
0.9298245614035088
0.9298245614035088
0.9298245614035088
0.9298245614035088
0.9298245614035088
0.9298245614035088
0.9298245614035088
0.9298245614035088
0.9385964912280702
0.9122807017543859
0.8859649122807017
```

## PROBLEM-2

```
import numpy as np # Import the NumPy library for numerical computing
import pandas as pd # Import the Pandas library for data manipulation and
analysis
import matplotlib.pyplot as plt # Import the Matplotlib library for data
visualization
from sklearn.preprocessing import StandardScaler # Import the StandardScaler
class from the scikit-learn library for data normalization

data = pd.read_csv("/content/sample_data/Housing.csv")

dft = data.drop(columns=['furnishingstatus']) # Create a new DataFrame "dft"
by dropping the 'furnishingstatus' column from the "data" DataFrame
col = dft.columns # Get the column names of the "dft" DataFrame and store
them in the "col" variable

from google.colab import drive # Import the 'drive' module from the
'google.colab' library for accessing Google Drive functionalities in Google
Colab
drive.mount('/content/drive') # Mount Google Drive to the '/content/drive'
directory in the Google Colab environment

Mounted at /content/drive
```



```
scaler = StandardScaler() # Create an instance of the StandardScaler class  
from the scikit-learn library for data normalization
```

```
#
```

```
def bin_map(var):  
    # Create a copy of the input variable to avoid modifying the original  
    data  
    var_copy = var.copy()
```

```
    # Loop through each element in the input variable
```

```
    for i in range(len(var_copy)):  
        # Convert 'yes' to 1 and 'no' to 0  
        if var_copy[i] == 'yes':  
            var_copy[i] = 1  
        elif var_copy[i] == 'no':  
            var_copy[i] = 0
```

```
    # Return the modified variable  
    return var_copy
```

```
# Define a mapping function to convert 'yes' to 1 and 'no' to 0
```

```
def bin_map(var):  
    # Create a copy of the input variable to avoid modifying the original  
    data  
    var_copy = var.copy()
```

```
    # Loop through each element in the input variable
```

```
    for i in range(len(var_copy)):  
        # Convert 'yes' to 1 and 'no' to 0  
        if var_copy[i] == 'yes':  
            var_copy[i] = 1  
        elif var_copy[i] == 'no':  
            var_copy[i] = 0
```

```
    # Return the modified variable  
    return var_copy
```

```
# Call the bin_map function on each column of the dft DataFrame
```

```
dft['mainroad'] = bin_map(dft['mainroad'])  
dft['guestroom'] = bin_map(dft['guestroom'])  
dft['basement'] = bin_map(dft['basement'])  
dft['hotwaterheating'] = bin_map(dft['hotwaterheating'])  
dft['airconditioning'] = bin_map(dft['airconditioning'])  
dft['prefarea'] = bin_map(dft['prefarea'])
```

```
dft = scaler.fit_transform(dft) # Use the `fit_transform` method of the  
`StandardScaler` object to normalize the data in `dft`
```

```
Y = scaler.fit_transform(np.array(data.price).reshape(545,1)) # Use the  
`fit_transform` method of the `StandardScaler` object to normalize the
```

*'price' column of the `data` DataFrame, reshape it to a column vector with 545 rows, and store the normalized data in the variable `Y`*

```
Wye = pd.DataFrame(Y) # Create a new DataFrame `Wye` from the normalized data `Y` using the `pd.DataFrame` constructor
```

```
data = pd.DataFrame(dft, columns=col) # Create a new DataFrame `data` from the normalized data `dft` using the `pd.DataFrame` constructor, specifying the column names as `col`
```

```
data # Display the `data` DataFrame
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom
\							
0	4.566365	1.046726	1.403419	1.421812	1.378217	0.405623	-0.465315
1	4.004484	1.757010	1.403419	5.405809	2.532024	0.405623	-0.465315
2	4.004484	2.218232	0.047278	1.421812	0.224410	0.405623	-0.465315
3	3.985755	1.083624	1.403419	1.421812	0.224410	0.405623	-0.465315
4	3.554979	1.046726	1.403419	-0.570187	0.224410	0.405623	2.149083
..	...	...	...	...	...	...	...
540	-1.576868	-0.991879	-1.308863	-0.570187	-0.929397	0.405623	-0.465315
541	-1.605149	-1.268613	0.047278	-0.570187	-0.929397	-2.465344	-0.465315
542	-1.614327	-0.705921	-1.308863	-0.570187	-0.929397	0.405623	-0.465315
543	-1.614327	-1.033389	0.047278	-0.570187	-0.929397	-2.465344	-0.465315
544	-1.614327	-0.599839	0.047278	-0.570187	0.224410	0.405623	-0.465315
	basement	hotwaterheating	airconditioning	parking	prefarea		
0	-0.734539	-0.219265	1.472618	1.517692	1.804941		
1	-0.734539	-0.219265	1.472618	2.679409	-0.554035		
2	1.361397	-0.219265	-0.679063	1.517692	1.804941		
3	1.361397	-0.219265	1.472618	2.679409	1.804941		
4	1.361397	-0.219265	1.472618	1.517692	-0.554035		
..	...	...	...	...	...		
540	1.361397	-0.219265	-0.679063	1.517692	-0.554035		
541	-0.734539	-0.219265	-0.679063	-0.805741	-0.554035		
542	-0.734539	-0.219265	-0.679063	-0.805741	-0.554035		
543	-0.734539	-0.219265	-0.679063	-0.805741	-0.554035		
544	-0.734539	-0.219265	-0.679063	-0.805741	-0.554035		

```
[545 rows x 12 columns]
```

```
Wye = pd.DataFrame(np.array(data.price)) # Create a new DataFrame `Wye` from the 'price' column of the `data` DataFrame using the `pd.DataFrame` constructor
```

```
Wye = pd.DataFrame(np.array(data.price)) # Create a new DataFrame `Wye` from the 'price' column of the `data` DataFrame using the `pd.DataFrame` constructor
```

```
train = data.sample(frac=0.8, random_state=1) # Create a new DataFrame
`train` by randomly sampling 80% of the rows from the `data` DataFrame using
the `sample` method, with a random seed of 1
```

```
test = data.drop(train.index) # Create a new DataFrame `test` by dropping
the rows from the `data` DataFrame that are present in the `train` DataFrame
using the `drop` method and passing the indices of `train` DataFrame as
argument
```

```
y_train = pd.DataFrame(np.array(train.price)) # Create a new DataFrame
`y_train` from the `price` column of the `train` DataFrame, converting it to
a NumPy array and then to a DataFrame
```

```
x_train = train.drop(columns=['price']) # Create a new DataFrame `x_train`
by dropping the `price` column from the `train` DataFrame using the `drop`
method
```

```
y_test = pd.DataFrame(np.array(test.price)) # Create a new DataFrame
`y_test` from the `price` column of the `test` DataFrame, converting it to a
NumPy array and then to a DataFrame
```

```
x_test = np.array(test.drop(columns=['price'])) # Create a NumPy array
`x_test` by dropping the `price` column from the `test` DataFrame using the
`drop` method, without converting it to a DataFrame
```

```
from sklearn.svm import SVR # Import the Support Vector Regression (SVR)
class from the scikit-Learn library
```

```
svr_rbf = SVR(kernel='rbf', C=1e3, gamma=0.1) # Create an instance of SVR
with RBF kernel, and set the hyperparameters C and gamma to 1e3 and 0.1,
respectively
```

```
svr_lin = SVR(kernel='linear', C=1e3) # Create an instance of SVR with
linear kernel, and set the hyperparameter C to 1e3
```

```
svr_poly = SVR(kernel='poly', C=1e3, degree=2) # Create an instance of SVR
with polynomial kernel of degree 2, and set the hyperparameters C and degree
to 1e3 and 2, respectively
```

```
y_rbf = svr_rbf.fit(x_train, y_train).predict(x_test) # Fit the SVR model
with RBF kernel using the training data (x_train, y_train), and predict the
target values for the test data (x_test). Store the predicted values in
y_rbf.
```

```
y_lin = svr_lin.fit(x_train, y_train).predict(x_test) # Fit the SVR model
with linear kernel using the training data (x_train, y_train), and predict
the target values for the test data (x_test). Store the predicted values in
y_lin.
```

```
y_poly = svr_poly.fit(x_train, y_train).predict(x_test) # Fit the SVR model
with polynomial kernel using the training data (x_train, y_train), and
predict the target values for the
```

```
/usr/local/lib/python3.9/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.9/dist-packages/sklearn/base.py:439: UserWarning: X
does not have valid feature names, but SVR was fitted with feature names
warnings.warn(
```

```
/usr/local/lib/python3.9/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.9/dist-packages/sklearn/base.py:439: UserWarning: X
does not have valid feature names, but SVR was fitted with feature names
warnings.warn(
```

```
/usr/local/lib/python3.9/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.9/dist-packages/sklearn/base.py:439: UserWarning: X
does not have valid feature names, but SVR was fitted with feature names
warnings.warn(
```

```
plt.plot(y_rbf) # Plot the predicted target values obtained from SVR model
with RBF kernel (y_rbf) as a line plot.
```

```
plt.plot(y_lin) # Plot the predicted target values obtained from SVR model
with linear kernel (y_lin) as a line plot.
```


```
plt.plot(y_poly) # Plot the predicted target values obtained from SVR model
with polynomial kernel (y_poly) as a line plot.
```

```
plt.xlabel('X') # Set the Label for x-axis as 'X'.
```

```
plt.ylabel('Y') # Set the Label for y-axis as 'Y'.
```

```
plt.title('SVR') # Set the title of the plot as 'SVR', indicating that it
represents the results of Support Vector Regression (SVR).
```

```
Text(0.5, 1.0, 'SVR')
```



```

import numpy as np # Import the NumPy library for numerical computing
import pandas as pd # Import the Pandas library for data manipulation and
analysis
import matplotlib.pyplot as plt # Import the Matplotlib library for data
visualization
from sklearn.preprocessing import StandardScaler # Import the StandardScaler
class from the scikit-learn library for data normalization

data = pd.read_csv("/content/sample_data/Housing.csv")

dft = data.drop(columns=['furnishingstatus']) # Create a new DataFrame "dft"
by dropping the 'furnishingstatus' column from the "data" DataFrame
col = dft.columns # Get the column names of the "dft" DataFrame and store
them in the "col" variable

scaler = StandardScaler() # Create an instance of the StandardScaler class
from the scikit-learn library for data normalization

#
def bin_map(var):
    # Create a copy of the input variable to avoid modifying the original
    data
    var_copy = var.copy()

    # Loop through each element in the input variable
    for i in range(len(var_copy)):

```

```

    # Convert 'yes' to 1 and 'no' to 0
    if var_copy[i] == 'yes':
        var_copy[i] = 1
    elif var_copy[i] == 'no':
        var_copy[i] = 0

# Return the modified variable
return var_copy

# Define a mapping function to convert 'yes' to 1 and 'no' to 0
def bin_map(var):
    # Create a copy of the input variable to avoid modifying the original
    data
    var_copy = var.copy()

    # Loop through each element in the input variable
    for i in range(len(var_copy)):
        # Convert 'yes' to 1 and 'no' to 0
        if var_copy[i] == 'yes':
            var_copy[i] = 1
        elif var_copy[i] == 'no':
            var_copy[i] = 0

    # Return the modified variable
    return var_copy

# Call the bin_map function on each column of the dft DataFrame
dft['mainroad'] = bin_map(dft['mainroad'])
dft['guestroom'] = bin_map(dft['guestroom'])
dft['basement'] = bin_map(dft['basement'])
dft['hotwaterheating'] = bin_map(dft['hotwaterheating'])
dft['airconditioning'] = bin_map(dft['airconditioning'])
dft['prefarea'] = bin_map(dft['prefarea'])

dft = scaler.fit_transform(dft) # Use the `fit_transform` method of the
`StandardScaler` object to normalize the data in `dft`

Y = scaler.fit_transform(np.array(data.price).reshape(545,1)) # Use the
`fit_transform` method of the `StandardScaler` object to normalize the
`price` column of the `data` DataFrame, reshape it to a column vector with
545 rows, and store the normalized data in the variable `Y`

Wye = pd.DataFrame(Y) # Create a new DataFrame `Wye` from the normalized
data `Y` using the `pd.DataFrame` constructor

data = pd.DataFrame(dft, columns=col) # Create a new DataFrame `data` from
the normalized data `dft` using the `pd.DataFrame` constructor, specifying
the column names as `col`

```

```
data # Display the `data` DataFrame
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom
0	4.566365	1.046726	1.403419	1.421812	1.378217	0.405623	-0.465315
1	4.004484	1.757010	1.403419	5.405809	2.532024	0.405623	-0.465315
2	4.004484	2.218232	0.047278	1.421812	0.224410	0.405623	-0.465315
3	3.985755	1.083624	1.403419	1.421812	0.224410	0.405623	-0.465315
4	3.554979	1.046726	1.403419	-0.570187	0.224410	0.405623	2.149083
...	...	...	...	...	...	...	...
540	-1.576868	-0.991879	-1.308863	-0.570187	-0.929397	0.405623	-0.465315
541	-1.605149	-1.268613	0.047278	-0.570187	-0.929397	-2.465344	-0.465315
542	-1.614327	-0.705921	-1.308863	-0.570187	-0.929397	0.405623	-0.465315
543	-1.614327	-1.033389	0.047278	-0.570187	-0.929397	-2.465344	-0.465315
544	-1.614327	-0.599839	0.047278	-0.570187	0.224410	0.405623	-0.465315

	basement	hotwaterheating	airconditioning	parking	prefarea
0	-0.734539	-0.219265	1.472618	1.517692	1.804941
1	-0.734539	-0.219265	1.472618	2.679409	-0.554035
2	1.361397	-0.219265	-0.679063	1.517692	1.804941
3	1.361397	-0.219265	1.472618	2.679409	1.804941
4	1.361397	-0.219265	1.472618	1.517692	-0.554035
...	...	...	...	...	...
540	1.361397	-0.219265	-0.679063	1.517692	-0.554035
541	-0.734539	-0.219265	-0.679063	-0.805741	-0.554035
542	-0.734539	-0.219265	-0.679063	-0.805741	-0.554035
543	-0.734539	-0.219265	-0.679063	-0.805741	-0.554035
544	-0.734539	-0.219265	-0.679063	-0.805741	-0.554035

```
[545 rows x 12 columns]
```

```
Wye = pd.DataFrame(np.array(data.price)) # Create a new DataFrame `Wye` from  
the 'price' column of the `data` DataFrame using the `pd.DataFrame`  
constructor
```

```
Wye = pd.DataFrame(np.array(data.price)) # Create a new DataFrame `Wye` from  
the 'price' column of the `data` DataFrame using the `pd.DataFrame`  
constructor
```

```
#d = pd.DataFrame(np.hstack([Ex,Wye]))
```

```
from sklearn.decomposition import PCA  
pca = PCA(n_components=2)  
principalComponents = pca.fit_transform(Ex)  
principalDf = pd.DataFrame(data = principalComponents  
    , columns = ['principal component 1', 'principal component 2'])
```

```
df1 = pd.DataFrame(np.hstack([principalDf, Wye])) # Creates a new DataFrame
by horizontally stacking two arrays: principalDf and Wye.
```

```
df1
```

	0	1	2
0	3.264248	-1.129485	4.566365
1	5.194952	-3.347516	4.004484
2	2.460935	1.278579	4.004484
3	3.625400	0.538743	3.985755
4	2.502535	1.070341	3.554979
..	...	...	...
540	-1.078671	1.444086	-1.576868
541	-2.533313	-0.848245	-1.605149
542	-2.063004	0.305558	-1.614327
543	-2.441185	-0.810540	-1.614327
544	-1.140103	-0.688090	-1.614327

```
[545 rows x 3 columns]
```

```
train = data.sample(frac=0.8, random_state=1) # Create a new DataFrame
`train` by randomly sampling 80% of the rows from the `data` DataFrame using
the `sample` method, with a random seed of 1
```

```
test = data.drop(train.index) # Create a new DataFrame `test` by dropping
the rows from the `data` DataFrame that are present in the `train` DataFrame
using the `drop` method and passing the indices of `train` DataFrame as
argument
```

```
from sklearn.svm import SVR # Importing Support Vector Regression (SVR)
from scikit-learn library
svr_rbf = SVR(kernel='rbf', C=1e3, gamma=0.1) # Creating an instance of SVR
with 'rbf' kernel, regularization parameter (C) set to 1e3, and gamma
parameter set to 0.1
```

```
train1ex = np.array(train1.drop(columns=[2])) # Extracting the data from
`train1` DataFrame by dropping the column with index 2 and converting it to a
numpy array
train1ex = train1ex.reshape(436, 2) # Reshaping the extracted data into a
2-dimensional array with 436 rows and 2 columns
```

```
train1we = np.array(train1.drop(columns=[0,1])) # Extracting the data from
`train1` DataFrame by dropping the columns with index 0 and 1, and converting
it to a numpy array
train1we = train1we.reshape(436, 1) # Reshaping the extracted data into a
2-dimensional array with 436 rows and 1 column
train1we = train1we.ravel() # Flattening the 2-dimensional array into a
1-dimensional array
```

```
test1ex = np.array(test1.drop(columns=[2])) # Extracting the data from
`test1` DataFrame by dropping the column with index 2, and converting it to a
```



*numpy array*

```
testlex = testlex.reshape(109, 2)  # Reshaping the extracted data into a  
2-dimensional array with 109 rows and 2 columns
```

```
testlwe = np.array(test1.drop(columns=[0,1]))  # Extracting the data from  
'test1' DataFrame by dropping the columns with index 0 and 1, and converting  
it to a numpy array
```

```
testlwe = testlwe.reshape(109, 1)  # Reshaping the extracted data into a  
2-dimensional array with 109 rows and 1 column
```

```
model = svr_rbf.fit(trainlex, trainlwe)  # Fitting the Support Vector  
Regression (SVR) model with the training data 'trainlex' and 'trainlwe'  
using RBF kernel
```

```
prediction1 = model.predict(testlex)  # Predicting the target values using  
the trained SVR model and test data 'testlex'
```


```
prediction1.reshape(109, 1)  # Reshaping the predicted values to  
have dimensions (109, 1)
```

```
array([[ 1.19151924e+00],  
       [ 3.30102310e+00],  
       [ 1.39622186e+00],  
       [ 3.93459752e-01],  
       [ 3.36186921e-01],  
       [ 4.36029068e+00],  
       [ 2.58189216e+00],  
       [ 1.96121431e+00],  
       [ 1.38799113e+00],  
       [ 5.98075163e-01],  
       [ 1.53558191e+00],  
       [ 4.23350039e-01],  
       [ 1.23999285e+00],  
       [ 6.28406224e-01],  
       [-2.84284115e-01],  
       [-9.25794346e-02],  
       [ 1.13813357e+00],  
       [ 1.14881210e+00],  
       [ 9.57050121e-01],  
       [ 2.11094127e-01],  
       [-4.59462076e-01],  
       [ 4.89055118e-01],  
       [ 9.88438638e-01],  
       [ 2.26987599e-01],  
       [ 1.18487895e+00],  
       [ 1.17244703e+00],  
       [ 1.39474671e+00],  
       [ 2.28369243e-01],  
       [ 7.05032576e-01],  
       [ 4.05263991e-01],  
       [ 9.53412617e-01],  
       [ 2.67012193e-01],
```

[ 3.18465224e-01],  
[ 3.63766240e-01],  
[ 8.22664454e-01],  
[-1.23963110e-01],  
[ 8.78462158e-01],  
[ 1.28957501e+00],  
[-1.34580596e-01],  
[-5.37043563e-02],  
[-3.04081703e-01],  
[-6.07783029e-01],  
[-4.67106739e-01],  
[-1.08885329e-01],  
[ 1.47711129e-01],  
[ 4.50823183e-01],  
[-4.29157181e-01],  
[-6.36425238e-01],  
[-3.80772185e-01],  
[-6.98067381e-01],  
[ 4.16601081e-01],  
[-2.96608270e-01],  
[ 1.54450890e-01],  
[-9.82117842e-01],  
[-6.92335414e-01],  
[-4.62576470e-01],  
[-2.20182045e-01],  
[-8.08757955e-01],  
[-8.08917695e-01],  
[-6.51788685e-01],  
[-3.80032844e-01],  
[-3.78707046e-01],  
[-3.61730447e-03],  
[-9.55285148e-02],  
[-1.08031964e-01],  
[-3.66645425e-01],  
[-3.44994357e-01],  
[ 1.10838364e-01],  
[ 8.24190922e-01],  
[-3.82576758e-02],  
[ 1.45870343e-01],  
[-6.46586193e-01],  
[ 3.23557673e-01],  
[-8.33166551e-01],  
[ 2.31803642e-01],  
[-3.97594414e-01],  
[-7.21799292e-01],  
[-8.90278452e-01],  
[-5.73060951e-01],  
[-5.96010997e-01],  
[-1.00565473e+00],  
[-2.64599980e-01],

```
[-9.03087824e-01],  
[-6.07433305e-01],  
[-7.46155877e-01],  
[-4.95912815e-01],  
[-7.74922359e-01],  
[-5.00362464e-01],  
[ 2.75375395e-01],  
[-8.56080255e-01],  
[ 1.19702367e-01],  
[ 2.26249858e-01],  
[-2.41687655e-01],  
[-6.22696587e-01],  
[-7.24296237e-01],  
[-1.08497870e+00],  
[-9.15844398e-01],  
[-5.94345980e-01],  
[-8.31366891e-01],  
[-8.78676322e-01],  
[-8.61433149e-01],  
[-8.79978297e-01],  
[-8.35936086e-01],  
[-7.39650680e-01],  
[-8.66176984e-01],  
[-5.48547119e-01],  
[-7.63062077e-01],  
[-7.38190753e-01],  
[-9.74750391e-01]])
```

```
plt.plot(test1we)           # Plotting the actual target values from test  
data  
plt.plot(prediction1, '+')  # Plotting the predicted target values with '+'  
marker  
[<matplotlib.lines.Line2D at 0x231b82396a0>]
```



```
import sklearn.metrics as sm

#Print Mean Absolute Error
print("Mean absolute error =", round(sm.mean_absolute_error(test1wye,
prediction1), 2))

#Print Mean Squared Error
print("Mean squared error =", round(sm.mean_squared_error(test1wye,
prediction1), 2))

#Print Median Absolute Error
print("Median absolute error =", round(sm.median_absolute_error(test1wye,
prediction1), 2))

#Print Explained Variance Score
print("Explain variance score =", round(sm.explained_variance_score(test1wye,
prediction1), 2))

#Print R2 Score
print("R2 score =", round(sm.r2_score(test1wye, prediction1), 2))

Mean absolute error = 0.47
Mean squared error = 0.45
Median absolute error = 0.33
Explain variance score = 0.56
R2 score = 0.56

from sklearn.decomposition import PCA
```

```

#Instantiate PCA with 6 components
pca = PCA(n_components=6)

#Fit and transform the data using PCA
principalComponents = pca.fit_transform(Ex)

#Create a DataFrame to store the principal components
principalDf = pd.DataFrame(data=principalComponents, columns=['principal
component 1', 'principal component 2', 'principal component 3', 'principal
component 4', 'principal component 5', 'principal component 6'])

#Concatenate the principal components with Wye
df2 = pd.DataFrame(np.hstack([principalDf, Wye]))

train2 = df2.sample(frac=0.8, random_state=1) # Create a training set by
randomly sampling 80% of the data from df2
test2 = df2.drop(train2.index) # Create a test set by removing the samples in
the training set from df2
from sklearn.svm import SVR # Import the SVR class from scikit-Learn
svr_rbf = SVR(kernel='rbf', C=1e3, gamma=0.1) # Create an SVR model with RBF
kernel, C=1e3, and gamma=0.1
train2 # Display the training set, which is a subset of df2 after random
sampling

```

	0	1	2	3	4	5	6
62	1.668434	-1.577144	-0.350531	-0.284510	-0.738116	-0.047705	1.232537
247	2.043753	-1.676954	-1.799614	0.760626	-0.463055	0.400701	-0.115977
142	1.691797	-1.073583	-0.967244	0.813824	-0.525831	0.919381	0.445904
107	0.706713	0.801670	0.287002	-0.536019	1.818314	1.209268	0.726844
483	-0.640626	-0.483671	-0.736872	-0.194905	0.549685	-0.005113	-0.977528
..	...	...	...	...	...	...	...
359	-1.203946	-0.065272	-0.719406	0.272859	-0.190551	0.152505	-0.565482
36	1.669630	-1.038315	-0.498018	3.720000	2.610809	0.056067	1.753214
30	2.528118	-2.072387	-1.250339	-0.374623	-0.620325	-0.739087	1.944253
20	0.044159	0.059643	0.115610	4.398738	0.810404	-1.162575	2.131547
527	-2.767009	0.582098	1.861415	-0.099819	-1.104261	0.482675	-1.333386

```

[436 rows x 7 columns]

train2ex = np.array(train2.drop(columns=[6])).reshape(436,6) # Extract the
features from the training set by dropping the column with index '6' and
reshaping the data into a 2D array with shape (436,6)

train2wye =
np.array(train2.drop(columns=[0,1,2,3,4,5])).reshape(436,1).ravel() # Extract
the target variable from the training set by dropping columns with indices
'0', '1', '2', '3', '4', '5', reshaping the data into a 1D array with shape
(436,) and flattening it

```

```

test2ex = np.array(test2.drop(columns=[6])).reshape(109,6) # Extract the
features from the test set by dropping column with index '6', reshaping the
data into a 2D array with shape (109, 6)


test2we = np.array(test2.drop(columns=[0,1,2,3,4,5])).reshape(109,1) #
Extract the target variable from the test set by dropping columns with
indexes 0 to 5, reshaping the data into a 2D array with shape (109, 1)

model = svr_rbf.fit(train2ex,train2we) # Fit the Support Vector Regression
(SVR) model using the training data
prediction2 = model.predict(test2ex) # Make predictions on the test data
using the trained model
prediction2.reshape(109, 1) # Reshape the prediction array to have shape
(109, 1) for plotting

plt.plot(test2we) # Plot the actual values of the target variable from the
test set
plt.plot(prediction2,'+') # Plot the predicted values of the target variable
using the SVR model

[<matplotlib.lines.Line2D at 0x231b816ec40>]

```



```

import sklearn.metrics as sm

#Print Mean Absolute Error
print("Mean absolute error =", round(sm.mean_absolute_error(test1we,
prediction1), 2))

#Print Mean Squared Error
print("Mean squared error =", round(sm.mean_squared_error(test1we,

```

```

prediction1), 2))

#Print Median Absolute Error
print("Median absolute error =", round(sm.median_absolute_error(test1wye,
prediction1), 2))

#Print Explained Variance Score
print("Explain variance score =", round(sm.explained_variance_score(test1wye,
prediction1), 2))

#Print R2 Score
print("R2 score =", round(sm.r2_score(test1wye, prediction1), 2))

Mean absolute error = 0.63
Mean squared error = 0.94
Median absolute error = 0.44
Explain variance score = 0.1
R2 score = 0.09

```

```

from sklearn.decomposition import PCA # Import the Principal Component
Analysis (PCA) module
pca = PCA(n_components=11) # Initialize the PCA model with 11 principal
components
principalComponents = pca.fit_transform(Ex) # Fit the PCA model on the input
data
principalDf = pd.DataFrame(data = principalComponents,
columns = ['principal component 1', 'principal component 2', 'principal
component 3',
'principal component 4', 'principal component 5', 'principal component 6',
'principal component 7', 'principal component 8', 'principal component 9',
'principal component 10', 'principal component 11']) # Create a dataframe to
store the principal components

,'principal component 9','principal component 10','principal component 11'
df3 = pd.DataFrame(np.hstack([principalDf,Wye])) # Combine the principal
components dataframe with the target variable dataframe to create a new
dataframe for further processing or analysis.

train3 = df3.sample(frac=0.8, random_state=1) # Create a training set by
randomly sampling 80% of the data from df3 with a random state of 1 for
reproducibility
test3 = df3.drop(train3.index) # Create a test set by removing the samples in
the training set from df3

from sklearn.svm import SVR # Import the Support Vector Regression (SVR)
module
svr_rbf = SVR(kernel='linear', C=1e3, gamma=0.1) # Initialize the SVR model

```

*with a linear kernel, C value of 1e3, and gamma value of 0.1*

*train3 # Display the training set for further examination or analysis*

	0	1	2	3	4	5	6	\
62	1.668434	-1.577144	-0.350531	-0.284510	-0.738116	-0.047705	-0.121506	
247	2.043753	-1.676954	-1.799614	0.760626	-0.463055	0.400701	-0.852142	
142	1.691797	-1.073583	-0.967244	0.813824	-0.525831	0.919381	-1.288370	
107	0.706713	0.801670	0.287002	-0.536019	1.818314	1.209268	0.212905	
483	-0.640626	-0.483671	-0.736872	-0.194905	0.549685	-0.005113	-0.635587	
..	...	...	...	...	...	...	...	
359	-1.203946	-0.065272	-0.719406	0.272859	-0.190551	0.152505	-0.658521	
36	1.669630	-1.038315	-0.498018	3.720000	2.610809	0.056067	1.965541	
30	2.528118	-2.072387	-1.250339	-0.374623	-0.620325	-0.739087	0.054959	
20	0.044159	0.059643	0.115610	4.398738	0.810404	-1.162575	1.328682	
527	-2.767009	0.582098	1.861415	-0.099819	-1.104261	0.482675	0.698249	
	7	8	9	10	11			
62	-0.160084	0.522901	0.678605	-1.033756	1.232537			
247	-0.194329	-2.276169	-0.701348	1.500006	-0.115977			
142	1.552603	0.030389	1.088125	-0.051945	0.445904			
107	-0.145298	-0.293530	0.394505	1.622250	0.726844			
483	0.720711	-0.405571	0.746588	0.217113	-0.977528			
..	...	...	...	...	...			
359	-0.470938	-0.347374	-0.178396	-0.807956	-0.565482			
36	0.918476	0.903971	-0.430595	0.405776	1.753214			
30	-0.117463	0.502399	-0.585040	1.183018	1.944253			
20	-1.589544	-0.499869	0.278410	0.796768	2.131547			
527	-0.361066	0.474880	-0.202054	1.022571	-1.333386			

[436 rows x 12 columns]

```
train3ex = np.array(train3.drop(columns=[11])).reshape(436,11)
```

*#Convert the training set to a numpy array and drop the column with index 11, then reshape it into a 2-dimensional array with 436 rows and 11 columns.*

```
train3we =  
np.array(train3.drop(columns=[0,1,2,3,4,5,6,7,8,9,10])).reshape(436,1).ravel()  
train3we =  
np.array(train3.drop(columns=[0,1,2,3,4,5,6,7,8,9,10])).reshape(436,1).ravel()  
)
```

*#Convert the target variable of the training set to a numpy array and drop the columns with indices 0 to 10, then reshape it into a 1-dimensional array with 436 elements using the ravel() function.*

```
test3ex = np.array(test3.drop(columns=[11])).reshape(109,11)
```

*#Convert the features of the test set to a numpy array and drop the column*



*with index 11, then reshape it into a 2-dimensional array with 109 rows and 11 columns.*

```
test3wye =  
np.array(test3.drop(columns=[0,1,2,3,4,5,6,7,8,9,10])).reshape(109,1)
```

*#Convert the target variable of the test set to a numpy array and drop the columns with indices 0 to 10, then reshape it into a 2-dimensional array with 109 rows and 1 column.*

```
#Import the SVR class from sklearn.svm module  
from sklearn.svm import SVR
```

```
#Create an instance of SVR with linear kernel and specified hyperparameters  
svr_rbf = SVR(kernel='linear', C=1e3, gamma=0.1)
```

```
#Fit the SVR model to the training data using train3ex as input features and  
train3wye as target variable  
model = svr_rbf.fit(train3ex,train3wye)
```


```
#Make predictions on the test data using the trained SVR model  
prediction3 = model.predict(test3ex)
```

```
#Reshape the predicted values to a 2D array with shape (109, 1)  
prediction3 = prediction3.reshape(109, 1)
```

```
#Plot the actual and predicted values  
plt.plot(test3wye, label='Actual')  
plt.plot(prediction3, label='Predicted')
```

```
#Add a Legend to the plot for identifying actual and predicted lines  
plt.legend(loc='best')
```

```
[<matplotlib.lines.Line2D at 0x231b7e30e20>]
```



```
import sklearn.metrics as sm

#Print Mean Absolute Error
print("Mean absolute error =", round(sm.mean_absolute_error(test1we,
prediction1), 2))

#Print Mean Squared Error
print("Mean squared error =", round(sm.mean_squared_error(test1we,
prediction1), 2))

#Print Median Absolute Error
print("Median absolute error =", round(sm.median_absolute_error(test1we,
prediction1), 2))

#Print Explained Variance Score
print("Explain variance score =", round(sm.explained_variance_score(test1we,
prediction1), 2))

#Print R2 Score
print("R2 score =", round(sm.r2_score(test1we, prediction1), 2))

Mean absolute error = 0.41
Mean squared error = 0.37
Median absolute error = 0.28
Explain variance score = 0.65
R2 score = 0.64

#Import the LinearRegression class from sklearn.linear_model module
from sklearn.linear_model import LinearRegression
```

```
#Create an instance of LinearRegression
```

```
model = LinearRegression()
```

```
#Create a dataframe "datalin" by horizontally stacking the "Ex" features and  
"Wye" target variable using np.hstack
```

```
datalin = pd.DataFrame(np.hstack([Ex, Wye]))
```

	0	1	2	3	4	5	6	\
0	1.046726	1.403419	1.421812	1.378217	0.405623	-0.465315	-0.734539	
1	1.757010	1.403419	5.405809	2.532024	0.405623	-0.465315	-0.734539	
2	2.218232	0.047278	1.421812	0.224410	0.405623	-0.465315	1.361397	
3	1.083624	1.403419	1.421812	0.224410	0.405623	-0.465315	1.361397	
4	1.046726	1.403419	-0.570187	0.224410	0.405623	2.149083	1.361397	
..	...	...	...	...	...	...	...	
540	-0.991879	-1.308863	-0.570187	-0.929397	0.405623	-0.465315	1.361397	
541	-1.268613	0.047278	-0.570187	-0.929397	-2.465344	-0.465315	-0.734539	
542	-0.705921	-1.308863	-0.570187	-0.929397	0.405623	-0.465315	-0.734539	
543	-1.033389	0.047278	-0.570187	-0.929397	-2.465344	-0.465315	-0.734539	
544	-0.599839	0.047278	-0.570187	0.224410	0.405623	-0.465315	-0.734539	
	7	8	9	10	11			
0	-0.219265	1.472618	1.517692	1.804941	4.566365			
1	-0.219265	1.472618	2.679409	-0.554035	4.004484			
2	-0.219265	-0.679063	1.517692	1.804941	4.004484			
3	-0.219265	1.472618	2.679409	1.804941	3.985755			
4	-0.219265	1.472618	1.517692	-0.554035	3.554979			
..	...	...	...	...	...			
540	-0.219265	-0.679063	1.517692	-0.554035	-1.576868			
541	-0.219265	-0.679063	-0.805741	-0.554035	-1.605149			
542	-0.219265	-0.679063	-0.805741	-0.554035	-1.614327			
543	-0.219265	-0.679063	-0.805741	-0.554035	-1.614327			
544	-0.219265	-0.679063	-0.805741	-0.554035	-1.614327			

```
[545 rows x 12 columns]
```

```
#Randomly sample 80% of the data for training
```

```
train_lin = datalin.sample(frac=0.8, random_state=1)
```

```
#Use the remaining data as testing set
```

```
test_lin = datalin.drop(train_lin.index)
```

```
#Get the shape of the testing set
```

```
test_lin_shape = test_lin.shape
```

```
(109, 12)
```

```
testlinY = test_lin.drop(columns=[0,1,2,3,4,5,6,7,8,9,10]) #Drop columns 0 to  
10 (inclusive) from the testing set to get the target variable
```

*#Train the linear regression model using the training data, drop column 11 from the training set as input features (X), drop columns 0 to 10 (inclusive) from the training set as target variable (y)*

```
model.fit(train_lin.drop(columns=[11]),train_lin.drop(columns=[0,1,2,3,4,5,6,7,8,9,10]))
```

```
LinearRegression()
```

*#Predict using the linear regression model*

```
pred = model.predict(test_lin.drop(columns=[11]))
```

```
import sklearn.metrics as sm
```

*#Print Mean Absolute Error*

```
print("Mean absolute error =", round(sm.mean_absolute_error(test1we, prediction1), 2))
```

*#Print Mean Squared Error*

```
print("Mean squared error =", round(sm.mean_squared_error(test1we, prediction1), 2))
```

*#Print Median Absolute Error*

```
print("Median absolute error =", round(sm.median_absolute_error(test1we, prediction1), 2))
```

*#Print Explained Variance Score*

```
print("Explain variance score =", round(sm.explained_variance_score(test1we, prediction1), 2))
```

*#Print R2 Score*

```
print("R2 score =", round(sm.r2_score(test1we, prediction1), 2))
```

Mean absolute error = 0.44

Mean squared error = 0.38

Median absolute error = 0.33

Explain variance score = 0.63

R2 score = 0.63